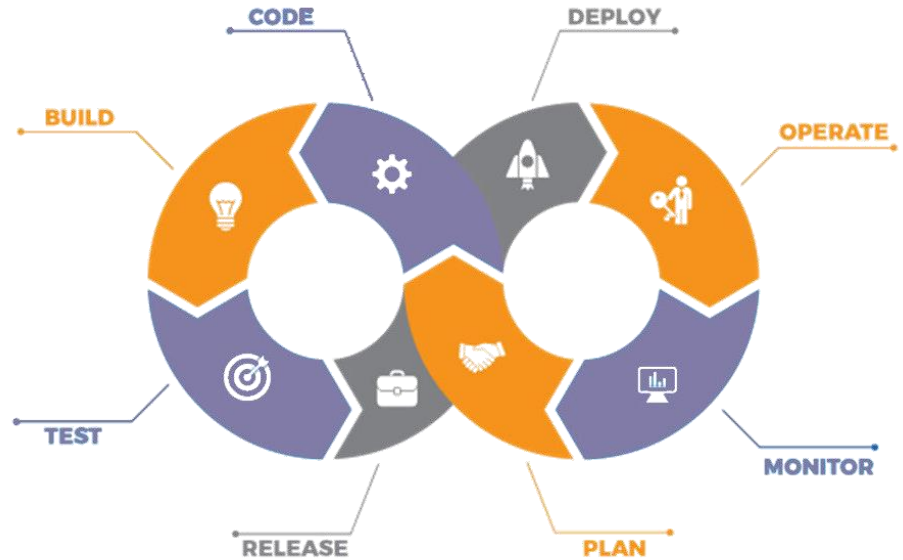


Introduction to Kubernetes



Agenda

01

Introduction to
Kubernetes

02

Docker Swarm Vs.
Kubernetes

03

Kubernetes
Architecture

04

Kubernetes
Installation

05

Working of
Kubernetes

06

Deployments in
Kubernetes

07

Services in
Kubernetes

08

Ingress in
Kubernetes

09

Kubernetes
Dashboard

Introduction to Kubernetes

Introduction to Kubernetes



- ★ Kubernetes is an open-source container orchestration software.
- ★ It was originally developed by Google.
- ★ It was first released on July 21, 2015.
- ★ It is the ninth most active repository on GitHub in terms of number of commits.


Features of Kubernetes

- ★ Pods
- ★ Replication Controller
- ★ Storage Management
- ★ Resource Monitoring
- ★ Health Checks
- ★ Service Discovery
- ★ Networking
- ★ Secret Management
- ★ Rolling Updates



Docker Swarm Vs. Kubernetes

Docker Swarm Vs. Kubernetes



Source: trends.google.com

Docker Swarm Vs. Kubernetes

Docker Swarm



- ★ Easy to install and initialize
- ★ Faster when compared to Kubernetes
- ★ Not reliable and has less features




Kubernetes

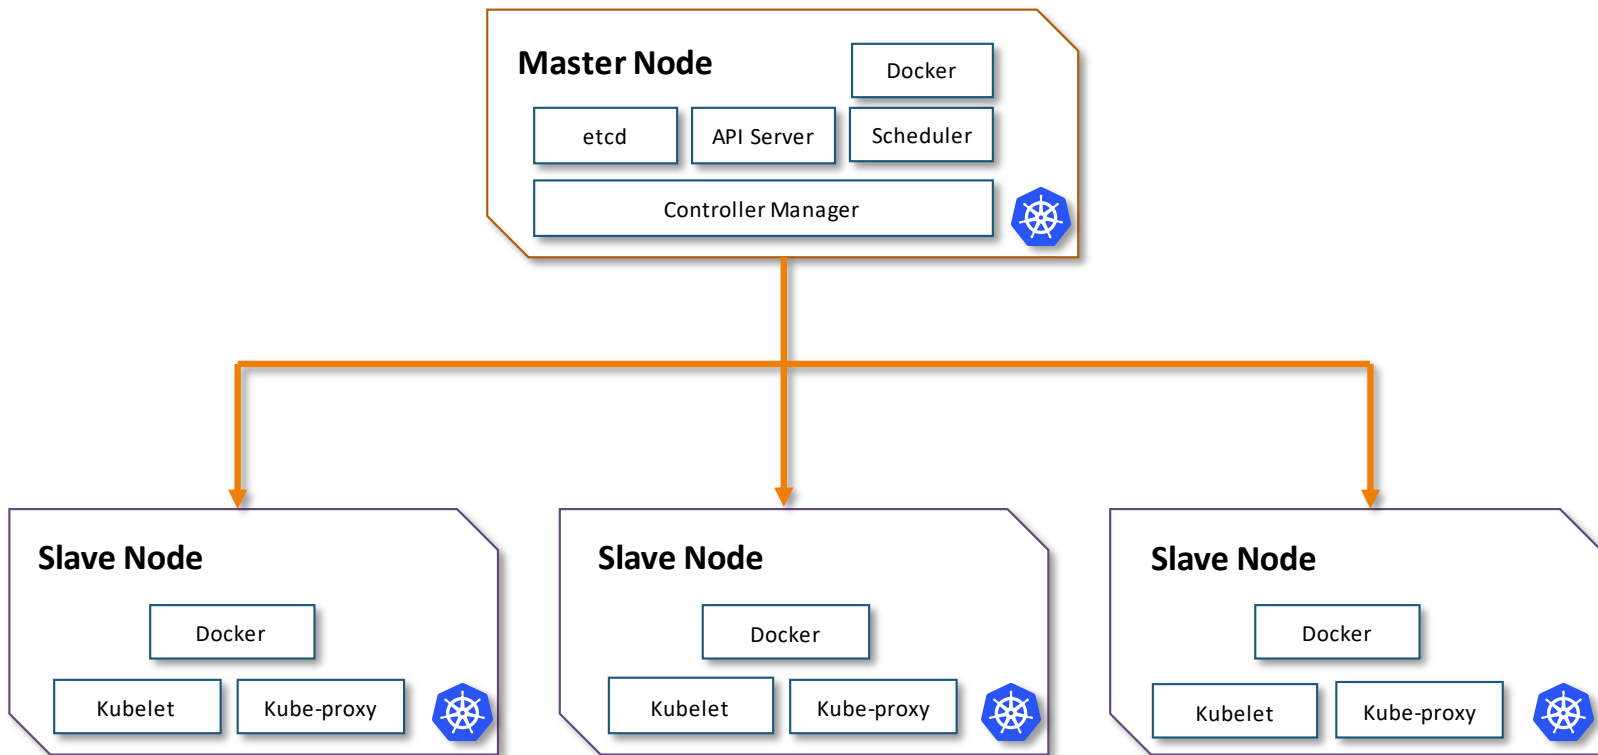
- ★ Complex procedure to install
- ★ Slower when compared to Docker Swarm
- ★ More reliable and has more features

Kubernetes Architecture

Kubernetes Architecture



Kubernetes Architecture



Kubernetes Architecture: Master Components

Kubernetes Architecture: Master Components

etcd

API Server

Scheduler

Controller Manager

It is a highly available distributed key–value store, which is used to store cluster wide secrets. It is only accessible by the Kubernetes API server, as it has sensitive information.

Master Node

etcd

API Server

Docker

Scheduler

Controller Manager



Kubernetes Architecture: Master Components

etcd

API Server

Scheduler

Controller Manager

It exposes Kubernetes API. Kubernetes API is the front-end for the Kubernetes Control Plane and is used to deploy and execute all operations in Kubernetes.

Master Node

etcd

API Server

Docker

Scheduler

Controller Manager



Kubernetes Architecture: Master Components

etcd

API Server

Scheduler

Controller Manager

The scheduler takes care of scheduling of all processes and the dynamic resource management and manages present and future events on the cluster.

Master Node

Docker

etcd


API Server

Scheduler

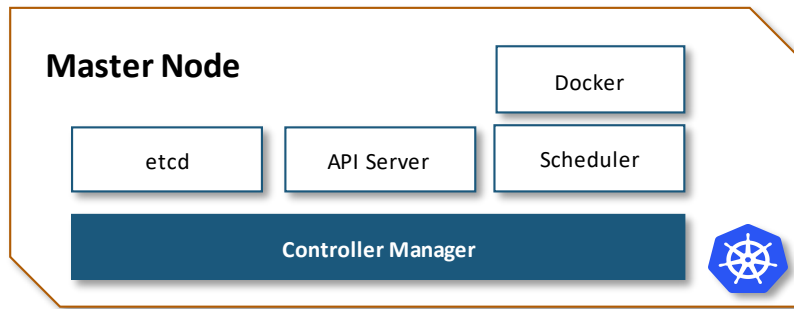
Controller Manager



Kubernetes Architecture: Master Components



The controller manager runs all controllers on the Kubernetes cluster. Although each controller is a separate process, to reduce complexity, all controllers are compiled into a single process. They are as follows:
Node Controller, Replication Controller, Endpoints Controller, Service Accounts and Token Controllers.



Kubernetes Architecture: Slave Components

Kubernetes Architecture: Slave Components

Kubelet

Kube-proxy

Kubelet takes the specification from the API server and ensures that the application is running according to the specifications which were mentioned. Each node has its own kubelet service.

Slave Node

Docker

Kubelet

Kube-proxy



Kubernetes Architecture: Slave Components

Kubelet

Kube-proxy

This proxy service runs on each node and helps in making services available to the external host. It helps in connection forwarding to the correct resources. It is also capable of doing primitive load balancing.

Slave Node

Docker

Kubelet

Kube-proxy




Kubernetes Installation

Kubernetes Installation

There are numerous ways to install Kubernetes. Following are some of the popular ways:

- **Kubeadm:** Bare Metal Installation
- **Minikube:** Virtualized Environment for Kubernetes
- **Kops:** Kubernetes on AWS
- **Kubernetes on GCP:** Kubernetes running on Google Cloud Platform



Hands-on: Installing Kubernetes Using Kubeadm

Working of Kubernetes

Working of Kubernetes



Pods can have one or more containers coupled together. They are the basic unit of Kubernetes. To increase high availability, we always prefer pods to be in replicas.



Pod – Replica 1




Pod – Replica 2



Pod – Replica 3

Working of Kubernetes

Services are used to load balance the traffic among the pods. It follows round-robin distribution among the healthy pods.



Pod – Replica 1

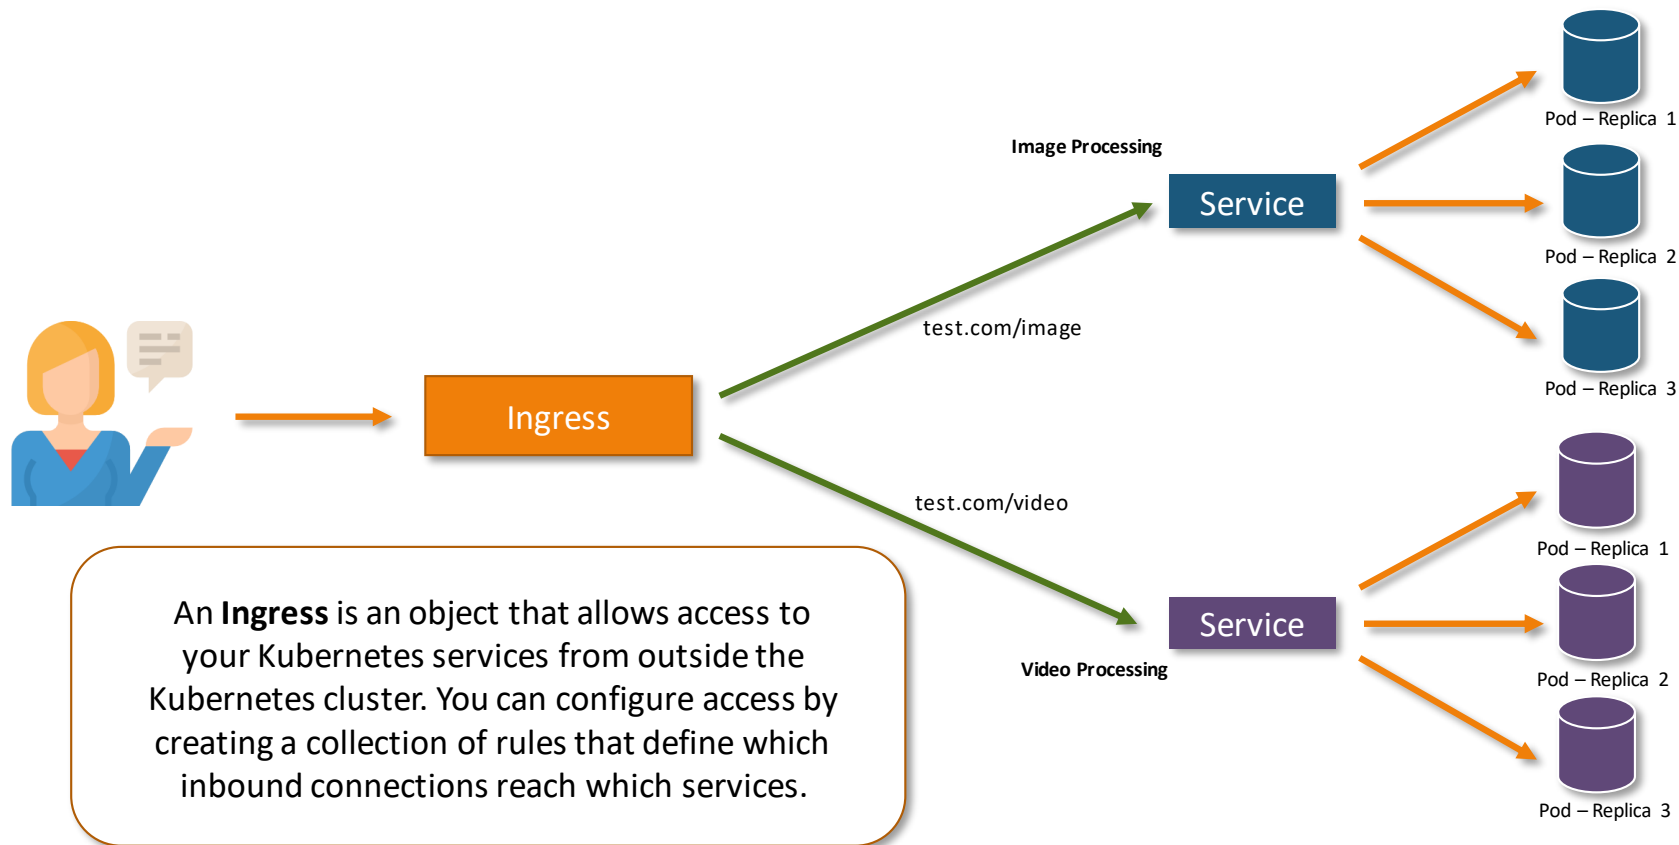


Pod – Replica 2



Pod – Replica 3

Working of Kubernetes




Deployments in Kubernetes

Deployments in Kubernetes

Deployment in Kubernetes is a controller which helps your applications reach the desired state; the desired state is defined inside the deployment file.

Deployment



Pods

YAML Syntax for Deployments

This YAML file will deploy 3 pods for nginx and will maintain the desired state, which is 3 pods, until this deployment is deleted.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
labels:
  app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.7.9
        ports:
        - containerPort: 80
```

Creating a Deployment

Once the file is created, to deploy this deployment use the following syntax:

Syntax

```
kubectl create -f nginx.yaml
```

 ubuntu@ip-172-31-39-244: ~

```
ubuntu@ip-172-31-39-244:~$ kubectl create -f nginx.yaml  
deployment.apps/nginx-deployment created  
ubuntu@ip-172-31-39-244:~$ █
```

Listing the Pods

To view the pods, type the following command:

Syntax

```
kubectl get po
```


```
ubuntu@ip-172-31-39-244: ~  
ubuntu@ip-172-31-39-244:~$ kubectl get po  
NAME                                READY   STATUS    RESTARTS   AGE  
nginx-deployment-76bf4969df-24vp1   1/1     Running   0           4m38s  
nginx-deployment-76bf4969df-frz7j   1/1     Running   0           4m38s  
nginx-deployment-76bf4969df-grnmc   1/1     Running   0           4m38s  
ubuntu@ip-172-31-39-244:~$
```

As you can see, the number of pods are matching with the number of replicas specified in the deployment file.

Creating a Service

Creating a Service

A Service is basically a round-robin load balancer for all pods, which matches with its name or selector. It constantly monitors the pods; in case a pod gets unhealthy, the service will start deploying the traffic to other healthy pods.




Service Types

ClusterIP: Exposes the service on cluster-internal IP

NodePort: Exposes the service on each Node's IP at a static port

LoadBalancer: Exposes the service externally using a cloud provider's load balancer

ExternalName: Maps the service to the DNS Name mentioned with the ExternalName service




Creating a NodePort Service

We can create a NodePort service using the following syntax:

Syntax

```
kubectl create service nodeport <name-of-service> --tcp=<port-of-service>:<port-of-container>
```

 ubuntu@ip-172-31-39-244: ~

```
ubuntu@ip-172-31-39-244:~$ kubectl create service nodeport nginx --tcp=80:80
service/nginx created
ubuntu@ip-172-31-39-244:~$ █
```

Creating a NodePort Service

To know the port, on which the service is being exposed, type the following command:

Syntax

```
kubectl get svc nginx
```

ubuntu@ip-172-31-39-244: ~

```
ubuntu@ip-172-31-39-244:~$ kubectl get svc nginx
```


NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
nginx	NodePort	10.103.235.81	<none>	80:32043/TCP	114s

```
ubuntu@ip-172-31-39-244:~$
```

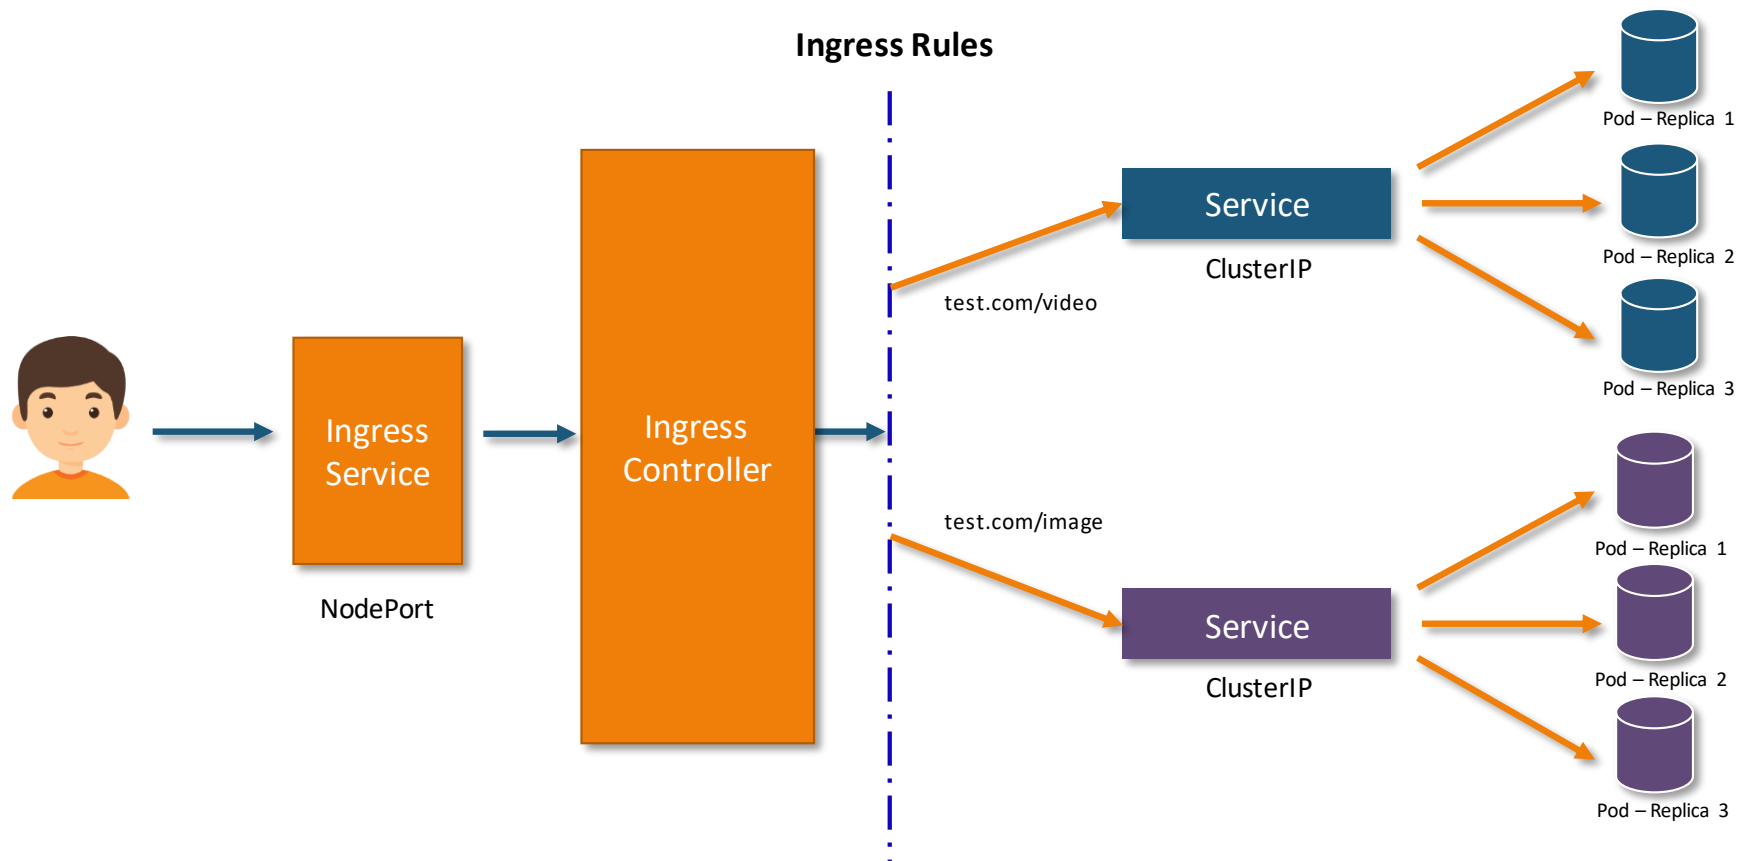
Creating an Ingress

What is an Ingress?

Kubernetes ingress is a collection of routing rules that govern how external users access services running in a Kubernetes cluster.



What is an Ingress?



Installing Ingress Controller

We will be using the nginx ingress controller for our demo. We can download it from the following link:

Link

<https://github.com/kubernetes/ingress-nginx/blob/master/docs/deploy/index.md>

The NGINX logo is displayed in a large, bold, green font. The letters are stylized, with the 'G' and 'i' having unique shapes. The 'X' is composed of two intersecting lines.

Defining Ingress Rules

The following rule, will redirect traffic which asks for /foo to nginx service. All other requests will be redirected to ingress controller's default page.

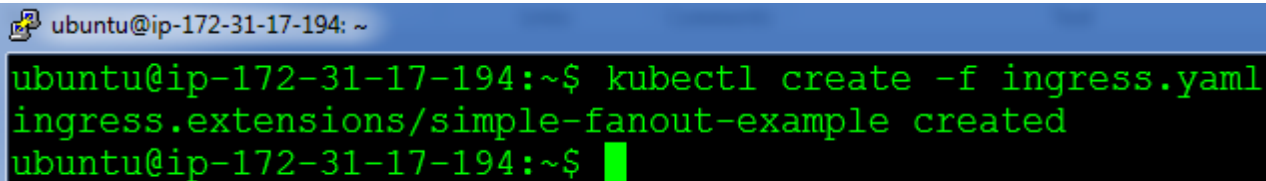
```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: simple-fanout-example
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
  - http:
      paths:
      - path: /foo
        backend:
          serviceName: nginx
          servicePort: 80
```

Deploying Ingress Rules

To deploy ingress rules, we use the following syntax:

Syntax

```
kubectl create -f ingress.yaml
```



A terminal window screenshot showing the execution of the `kubectl create -f ingress.yaml` command. The prompt is `ubuntu@ip-172-31-17-194: ~`. The output shows the command being executed and the message `ingress.extensions/simple-fanout-example created`. The prompt then returns to `ubuntu@ip-172-31-17-194:~$` with a green cursor.

```
ubuntu@ip-172-31-17-194: ~  
ubuntu@ip-172-31-17-194:~$ kubectl create -f ingress.yaml  
ingress.extensions/simple-fanout-example created  
ubuntu@ip-172-31-17-194:~$
```

Viewing Ingress Rules

To list the ingress rules we use the following syntax:


Syntax

```
kubectl get ing
```

```
ubuntu@ip-172-31-17-194: ~  
ubuntu@ip-172-31-17-194:~$ kubectl get ing  
NAME                HOSTS    ADDRESS    PORTS    AGE  
simple-fanout-example *        80         2m5s  
ubuntu@ip-172-31-17-194:~$
```

Kubernetes: Building Blocks: Pod


- Pod is the **smallest** and **simplest** Kubernetes **object**.
- It is the **unit of deployment** in Kubernetes, which represents a single instance of the application.
- A Pod is a **logical collection** of one or more **containers**, which:
 - Are **scheduled** together on the same host
 - Share the same **network** namespace
 - **Mount** the same external storage (**volumes**)
- Pods are **ephemeral** in nature, and they do not have the capability to **self-heal** by themselves.
- We use them with **controllers** like **Deployments**, **ReplicaSets**, which can handle a Pod's **replication**, **fault tolerance**, **self-heal**, etc



Kubernetes: Building Blocks: Labels

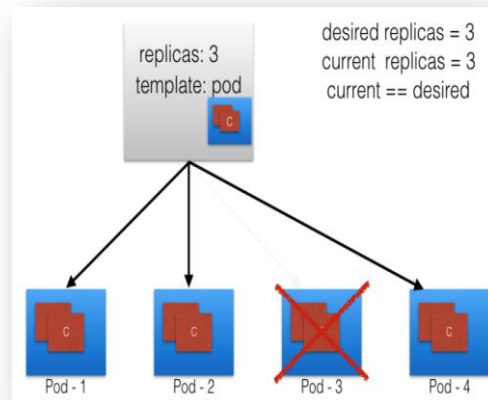
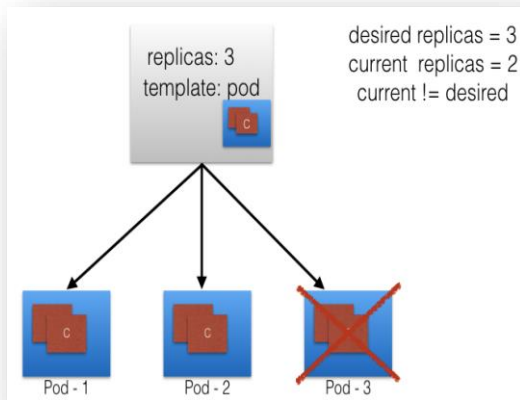
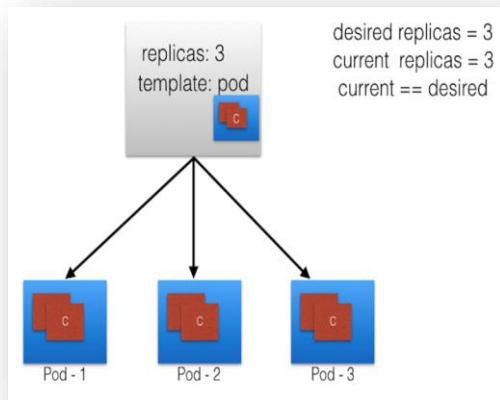
- **Labels** are **key-value** pairs that can be **attached** to any **Kubernetes objects** (e.g. Pods).
- Labels are used to **organize** and **select** a **subset** of **objects**, based on the **requirements** in place.
- Many objects can have the same Label(s).
- Labels do not provide **uniqueness** to objects.
- We have used two Labels: app and env.
- Based on our requirements, we have given different values to our four Pods

- **Label Selector**
- We can use Label Selectors, to select subset of objects
- For example,
- **Equity Based Selector**
env==dev, env!=prod ,
- **Set Based Selector**
env in (dev,qa), app notin (backend,middleware)



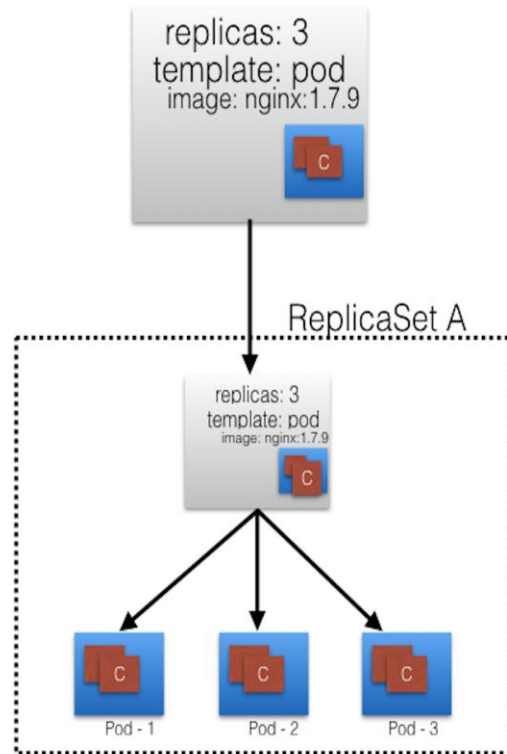
Kubernetes: Building Blocks: ReplicaSets

- **ReplicationController** is a **controller** that is part of the master node's **controller manager**.
- It makes sure the **specified** number of **replicas** for a Pod is **running** at any given point in time.
- If there are more Pods than the **desired** count, the ReplicationController would **kill** the extra **Pods**, and, if there are less Pods, then the ReplicationController would **create more Pods** to match the **desired** count.
- **ReplicaSet** is the **next-generation ReplicationController**.
- **ReplicaSets** support both equality- and set-based selectors, whereas **ReplicationControllers** only support equality-based Selectors.



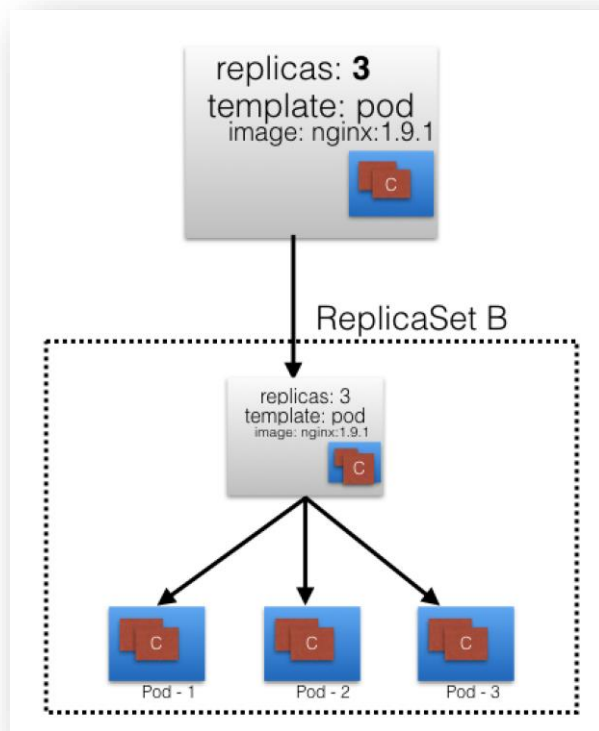
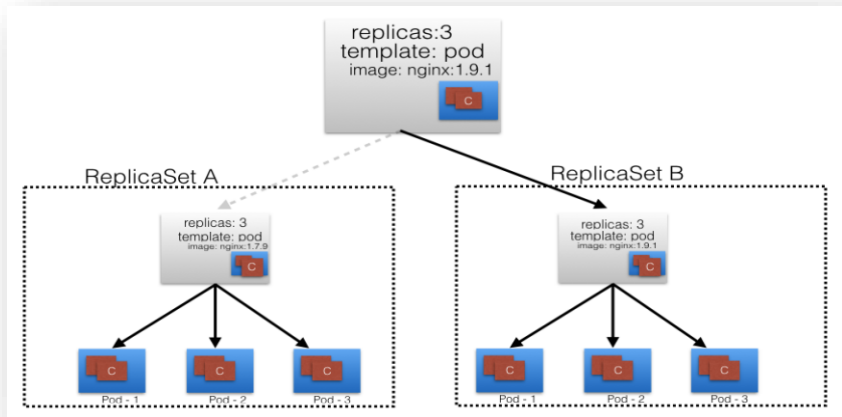
Kubernetes: Building Blocks: Deployments

- **Deployment** objects provide **declarative** updates to Pods and **ReplicaSets**.
- The **DeploymentController** is part of the master node's **controller manager**, and it makes sure that the current state always matches the **desired state**.
- In the following example, we have a **Deployment** which creates a **ReplicaSet A**.
- ReplicaSet A then creates 3 **Pods**.
- In each **Pod**, one of the **containers** uses the **nginx:1.7.9** image.



Kubernetes: Building Blocks: Deployment Rollout

- Now, in the **Deployment**, we change the **Pods Template** and we **update the image** for the nginx container from nginx:1.7.9 to nginx:1.9.1.
- As have **modified the Pods Template**, a new **ReplicaSet B** gets created.
- This process is referred to as a **Deployment rollout**.
- A **rollout** is only **triggered** when we **update the Pods Template** for a **deployment**.
- **Operations** like **scaling** the deployment **do not trigger** the deployment.
- Once **ReplicaSet B** is ready, the Deployment starts pointing to it.




Kubernetes Networking

- In Kubernetes, each Pod gets a **unique** IP address.
- Kubernetes uses **Container Network Interface** (CNI), proposed by **CoreOS** to assign the IP address to each **Pod**.
- Inside a Pod, **containers** share the **network namespaces**, so that they can reach to each other via **localhost**.
- **Pod-to-Pod** communication across **Nodes** can be achieved via:
 - **Routable** Pods and nodes, using the **underlying physical infrastructure**, like **Google Kubernetes Engine**
 - Using **Software Defined Networking**, like Flannel, Weave, Calico, etc.
- We can access our **applications** from outside the **cluster** by exposing our **services** to the external world with **kube-proxy**

Kubernetes: Volumes

- Containers inside pods are **ephemeral** in nature
- All data stored inside a container is **deleted** if the container **crashes**
- kubelet will **restart** it with a **clean state**, which means that it will not have any of the **old data**
- To overcome this problem, Kubernetes uses **Volumes**
- A Volume is essentially a **directory** backed by a **storage** medium.
- The storage medium and its content are determined by the Volume Type
- Volume is attached to a **Pod** and shared among the **containers** of that Pod
- The **Volume** has the same **life span** as the **Pod**, and it outlives the containers of the Pod - this allows **data** to be **preserved** across container restarts.

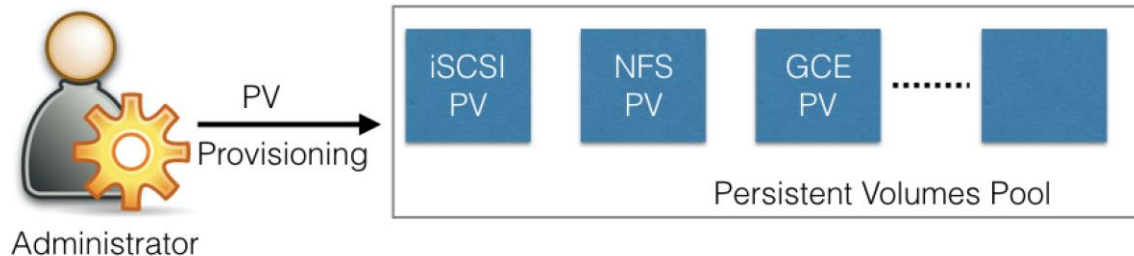


Kubernetes: Volumes Types

- A directory which is **mounted** inside a **Pod** is backed by the underlying **Volume Type**
- Volume Type decides the **properties** of the directory, like **size**, **content**, etc
- Some examples of Volume Types are:
 - emptyDir
 - hostPath
 - gcePersistentDisk
 - awsElasticBlockStore
 - nfs
 - iscsi
 - secret
 - persistentVolumeClaim

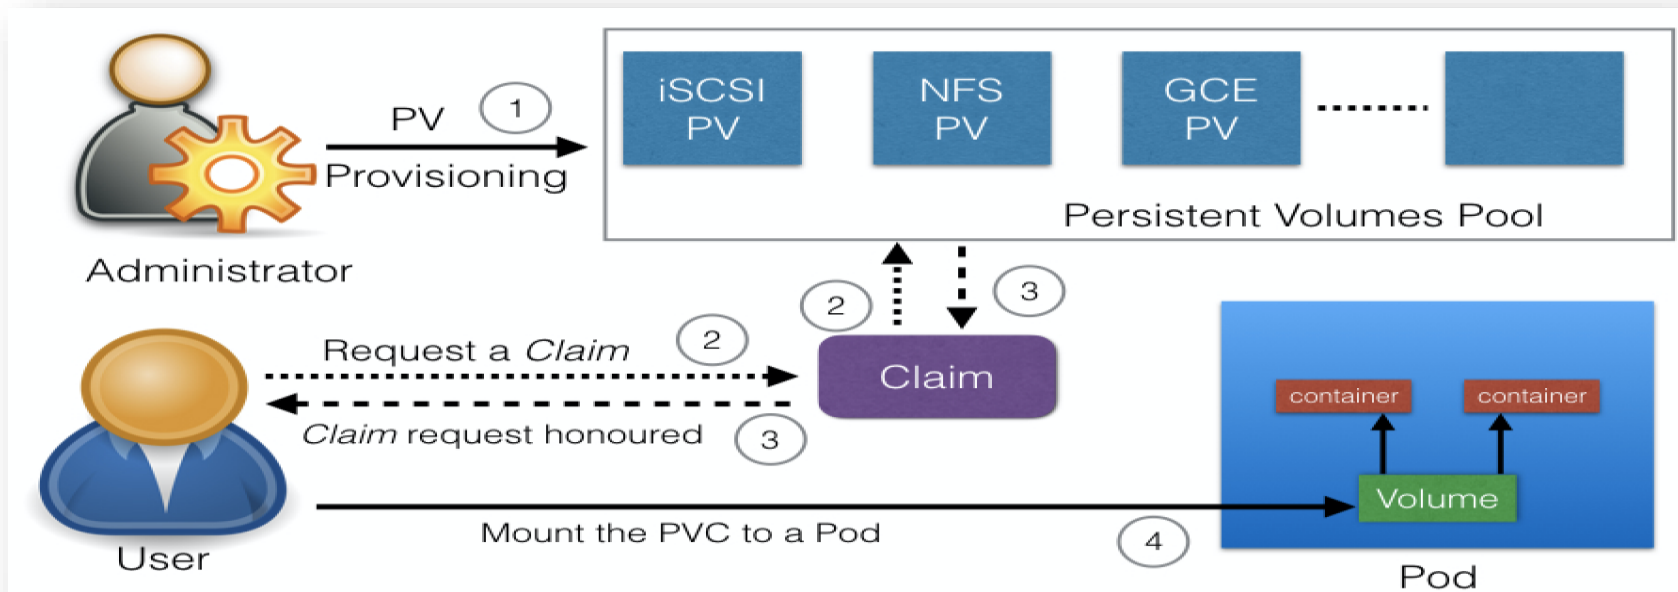
Kubernetes: PersistentVolumes

- **Persistent Volume** is a **network-attached storage** in the **cluster**, which is provisioned by the **administrator**
- PersistentVolumes can be **dynamically** provisioned based on the **StorageClass** resource
- A **StorageClass** contains pre-defined **provisioners** and parameters to create a **PersistentVolume**
- Using **PersistentVolumeClaims**, a user sends the request for dynamic PV **creation**, which gets wired to the **StorageClass** resource.
- Some of the Volume Types that support managing storage using PersistentVolumes are:
 - GCEPersistentDisk
 - AWSElasticBlockStore
 - AzureFile
 - NFS
 - iSCSI



Kubernetes: PersistentVolumeClaims

- **PersistentVolumeClaim (PVC)** is a request for **storage** by a **user**.
- Users **request** for **PersistentVolume** resources based on size, access modes, etc.
- Once a suitable **PersistentVolume** is found, it is bound to a **PersistentVolumeClaim**.
- Once a user finishes its work, the attached PersistentVolumes can be released.
- The underlying PersistentVolumes can then be reclaimed and recycled for future usage.



Kubernetes: Liveness Probe

- Liveness probe checks on an application's health, and, if for some reason, the health check fails, it restarts the affected container automatically
- Liveness Probes can be set by defining:
 - Liveness command
 - Liveness HTTP request
 - TCP Liveness Probe.

- **Liveness Command**

In this demo file `liveness-exec.yaml`, we will use existence of on file inside container as **Health Criteria**.

File: `/tmp/healthy`

Initial Probe: after 3 seconds

Probe Interval: 5 seconds

File deletion: 30 seconds after creation

- `sudo kubectl create -f liveness-exec.yaml`
- `sudo kubectl get pods`
- `sudo kubectl describe pod liveness-exec`

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    test: liveness
  name: liveness-exec
spec:
  containers:
    - name: liveness
      image: k8s.gcr.io/busybox
      args:
        - /bin/sh
        - -c
        - touch /tmp/healthy; sleep 30; rm -rf /tmp/healthy; sleep
600
      livenessProbe:
        exec:
          command:
            - cat
            - /tmp/healthy
          initialDelaySeconds: 3
          periodSeconds: 5
```

Kubernetes: Liveness Probe

- **Liveness HTTP**

In this demo file liveness-http.yaml ,
we will use one http response inside container
as **Health Criteria** .

- `sudo kubectl create -f liveness-http.yaml`
- `sudo kubectl get pods`
- `sudo kubectl describe pod liveness-http`

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    test: livenesshttp
  name: liveness-http
spec:
  containers:
  - name: livenesshttp
    image: httpd:latest
    ports:
      - containerPort: 80
    livenessProbe:
      httpGet:
        path: /healthz
        port: 80
        httpHeaders:
          - name: X-Custom-Header
            value: Awesome
      initialDelaySeconds: 3
      periodSeconds: 3
```

Kubernetes: hostPath Volume Type

- In this Demo we will Create Nginx Service which will use hostPath Volume to get Content
- Create a Directory named “vol” on Node Machine
- Create a file named “index.html” with given content
- Create a Deployment from webserver-vol.yaml
- Create a Service from webserver-svc-vol.yaml
- Describe Service to get nodePort
- Visit Node's IP_Address:nodePort in browser
- It will show you Nginx page with custom home page
- change index.html's content and refresh tab in browser

```
<html>
<head>K8S Volume Demo</head>
<body>
<h1>
Welcome to Kubernetes Session
</h1>
</body>
</html>
index.html
```

```
apiVersion: v1
kind: Service
metadata:
  name: web-service
  labels:
    run: web-service
spec:
  type: NodePort
  ports:
    - port: 80
      protocol: TCP
  selector:
    app: webserver
```

webserver-svc-vol.yaml

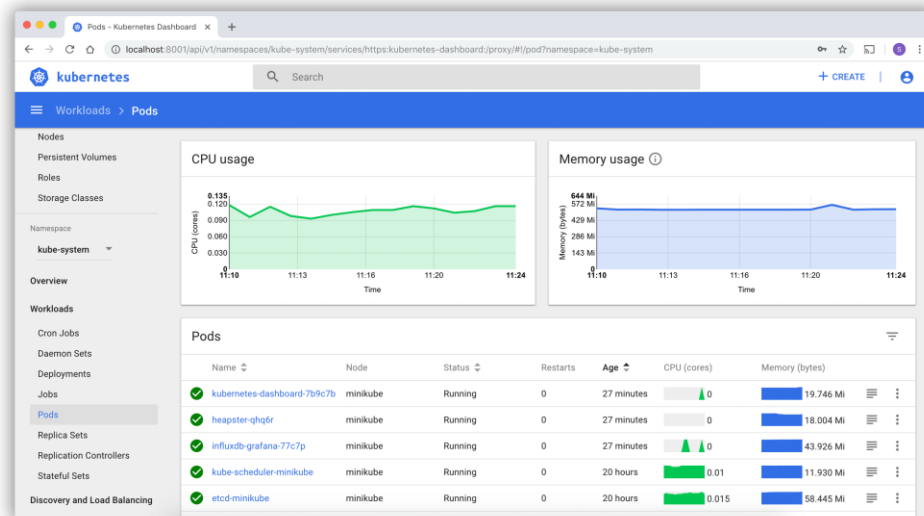
webserver-vol.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: webserver
spec:
  replicas: 1
  selector:
    matchLabels:
      app: webserver
  template:
    metadata:
      labels:
        app: webserver
    spec:
      containers:
        - name: webserver
          image: nginx:alpine
          ports:
            - containerPort: 80
          volumeMounts:
            - name: hostvol
              mountPath: /usr/share/nginx/html
          volumes:
            - name: hostvol
```


Kubernetes Dashboard

Kubernetes Dashboard

Dashboard is a web-based Kubernetes user interface. You can use Dashboard to deploy containerized applications to a Kubernetes cluster, troubleshoot your containerized application and manage cluster resources.

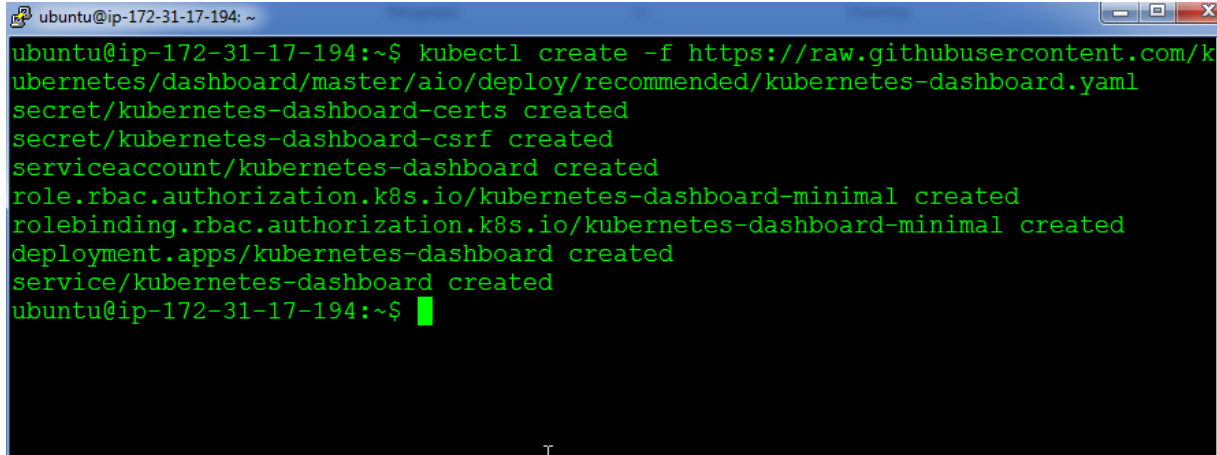


Installing Kubernetes Dashboard

To install Kubernetes Dashboard, execute the following command:

Syntax

```
kubectl create -f  
https://raw.githubusercontent.com/kubernetes/dashboard/v1.10.1/src/deploy/recommended/kubernetes-  
dashboard.yaml
```

A terminal window with a blue title bar showing the command to install the Kubernetes Dashboard. The command is executed successfully, and the output lists the created resources: secrets for certificates and CSRF, a service account, a role, a role binding, and the dashboard deployment and service.

```
ubuntu@ip-172-31-17-194: ~  
ubuntu@ip-172-31-17-194:~$ kubectl create -f https://raw.githubusercontent.com/kubernetes/dashboard/v1.10.1/src/deploy/recommended/kubernetes-dashboard.yaml  
secret/kubernetes-dashboard-certs created  
secret/kubernetes-dashboard-csrf created  
serviceaccount/kubernetes-dashboard created  
role.rbac.authorization.k8s.io/kubernetes-dashboard-minimal created  
rolebinding.rbac.authorization.k8s.io/kubernetes-dashboard-minimal created  
deployment.apps/kubernetes-dashboard created  
service/kubernetes-dashboard created  
ubuntu@ip-172-31-17-194:~$
```

Accessing Kubernetes Dashboard

Change the service type for kubernetes-dashboard to NodePort

Syntax

```
kubectl -n kubernetes-dashboard edit service kubernetes-dashboard
```

```
ubuntu@ip-172-31-17-194: ~  
[1] Please edit the object below. Lines beginning with a '#' will be ignored,  
# and an empty file will abort the edit. If an error occurs while saving,  
# it will be reopened with the relevant failures.  
#  
apiVersion: v1  
kind: Service  
metadata:  
  creationTimestamp: "2019-02-05T10:16:53Z"  
  labels:  
    k8s-app: kubernetes-dashboard  
  name: kubernetes-dashboard  
  namespace: kube-system  
  resourceVersion: "21192"  
  selfLink: /api/v1/namespaces/kube-system/services/kubernetes-dashboard  
  uid: 287flaa5-292f-11e9-ab4d-0689f8984fe2  
spec:  
  clusterIP: 10.104.60.164  
  externalTrafficPolicy: Cluster  
  ports:  
    - nodePort: 30788  
      port: 443  
      protocol: TCP  
      targetPort: 8443  
  selector:  
    k8s-app: kubernetes-dashboard  
  sessionAffinity: None  
  type: NodePort  
status:  
  loadBalancer: {}
```

Logging into Kubernetes Dashboard

1. Check the NodePort from the kubernetes-dashboard service
2. Browse to your cluster on the Internet browser, and enter the IP address
3. Click on Token, which will ask you for the token entry
4. Generate a token using the following command:

```
$ kubectl create serviceaccount cluster-admin-dashboard-sa
$ kubectl create clusterrolebinding cluster-admin-dashboard-sa \
  --clusterrole=cluster-admin \
  --serviceaccount=default:cluster-admin-dashboard-sa

$ TOKEN=$(kubectl describe secret $(kubectl -n kube-system get secret | awk '/^cluster-admin-dashboard-sa-token-/{print $1}') | awk '$1=="token:"{print $2}')

$ echo $TOKEN
```

5. Finally, enter the token and login to your dashboard

Hands-on: Deploying an App Using Dashboard